



Random Acts of Pizza

W207: Spring '17 Final Project
Erika Lawrence, Geoff Stirling, Jen Jen Chen, Leslie Teo

Project Overview

- Reddit community: Random Acts of Pizza
- Can ML help us predict who gets a free pizza?
- Data set of 5,671 text-based requests for cheesy goodness



Photo credit: thr.com, Hollywood Reporter, 2017

Solution Process & Tools

- The Data

- Text

- Request

- Title

- Metadata

- Account: reddit, history

- Post: age, votes

request_title

Request Colorado Springs Help Us Please
California, No cash and I could use ...
Hungry couple in Dundee, Scotland wo...


requester_subreddits_at_request

[FoodstuffsAllAround, IAmA, RandomActsOfCookie...
[AskReddit, Guitar, Jazz, Music, NSFW_GIF, Psy...
[Albany, AskReddit, Brooklyn, Favors, ImGoingT...
[AskReddit, Favors, Jazz, boardgames, discgolf...



Solution Process & Tools

- Data cleaning/pre-processing
 - Text
 - Normalize/Denormalize
 - Vectorize
 - Clean
 - Lemmatize ('reprocessing' -> 'process')



Solution Process & Tools

- Metadata features:
 - Time/date of request
 - Up/down votes
- Sentiment Analysis
 - vaderSentiment




Solution Process & Tools

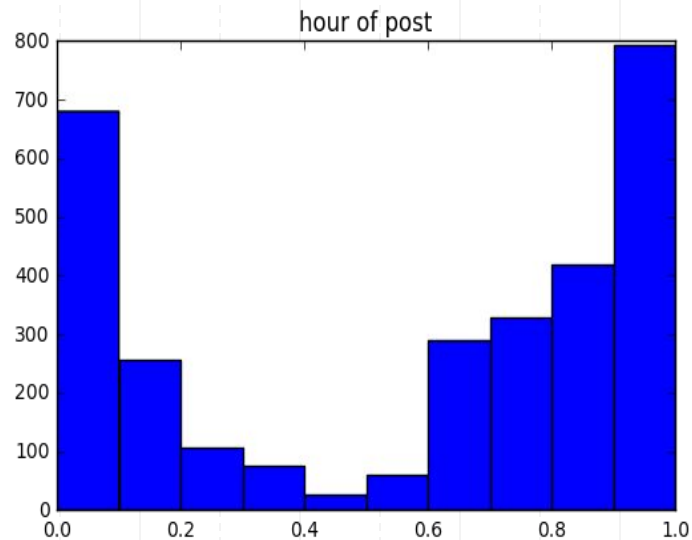
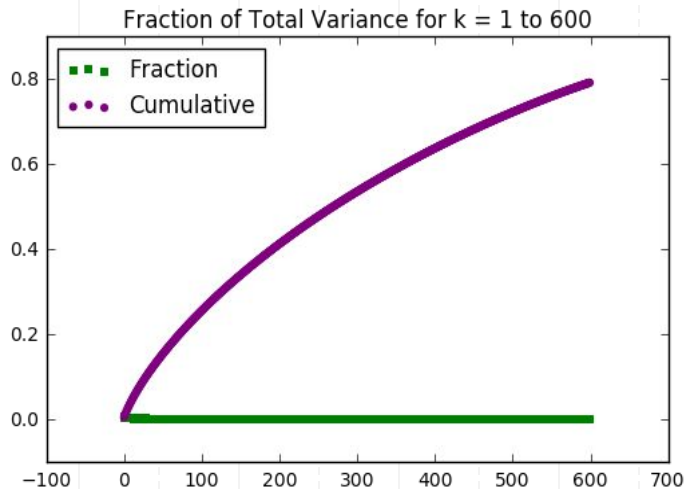
- Methods/Models:
 - Vectorize: bi/tri-gram, tf-idf
 - PCA
 - Logistic regression
 - Bernoulli NB
 - Random Forest
 - Ensemble
 - (+ SVM, Ada Forest, Gradient Boosting, and Decision Trees)



Results of our work: baseline 74%



Results of our work: features



Results of our work: final


Random Forest with n=100

Accuracy: 0.837

Area under curve of ROC: 0.721465696466

Classification report:

	precision	recall	f1-score	support
False	0.84	0.96	0.90	740
True	0.82	0.48	0.61	260
avg / total	0.83	0.84	0.82	1000



What we learned for next time

- Expand time analysis
- Vader expansion
- Features are key. Start early.
- Never give up. Never surrender.



THANKS!

Any questions?

